

HTTP/2

Kevin Scrivnor

COMP 429

Fall 2023

Review

- HTTP is a **request/response** model for delivering web applications
- HTTP is an **ASCII based protocol**

- Statistics (Spring 2023)

- Total kilobytes: median **2340**
- Total request: median **71**
- https request: **95%**
- TCP conn. Per page: median **12**
- HTTP/2 requests: **67.5%**
- (new stat) HTTP/3 support: **18.7%**
- <https://httparchive.org/reports/state-of-the-web>

Statistics (Spring 2021)

Total kilobytes: median 2038
Total request: median 73
https request: 89.3%
TCP conn. Per page: median 13
HTTP/2 requests: 66.9%

Statistics (Fall 2021)

Total kilobytes: median 2174
Total request: median 74
https request: 92.3%
TCP conn. Per page: median 14
HTTP/2 requests: 66.8%

Heart of the Internet

- **TCP/IP** was first proposed together in 1974 as the “Internet Protocol Suite”
 - *A solution for transmitting data across an unreliable channel.*
 - 1981, we have our two RFCs published
 - IPv4, RFC 791
 - TCP, RFC 793
- **TCP**
 - Complexity of dealing with the unreliable channel is abstracted away
 - Retransmission of lost data
 - In-order delivery (to the application)
 - Congestion control and avoidance
 - Data Integrity
 - Etc.
 - For now, we just need to know general concepts
 - *TCP does a thing, not necessarily TCP does a thing and this is how it works.*


TCP doesn't necessarily help modern day web browsing

- HTTP standard does not state to use TCP
- In practice, **all HTTP traffic is delivered over TCP**
 - Though potentially that will change with QUIC which is still reliable
- TCP makes guarantees:
 - In order delivery of data from the server to the client
 - Mechanisms to avoid flooding the channel with data
 - An established connection
- But all of this comes at a cost!

TCP Handshake

- New York to London connection establishment
- **SYN, SYNACK, ACK**
- Before the data is sent
- Connection establishment is expensive
- Critical for HTTP to *reuse connections* that are already established
- *Note the delay before the data is sent*

Route	Distance	Time, light in vacuum	Time, light in fiber	Round-trip time (RTT) in fiber
New York to San Francisco	4,148 km	14 ms	21 ms	42 ms
New York to London	5,585 km	19 ms	28 ms	56 ms
New York to Sydney	15,993 km	53 ms	80 ms	160 ms
Equatorial circumference	40,075 km	133.7 ms	200 ms	200 ms



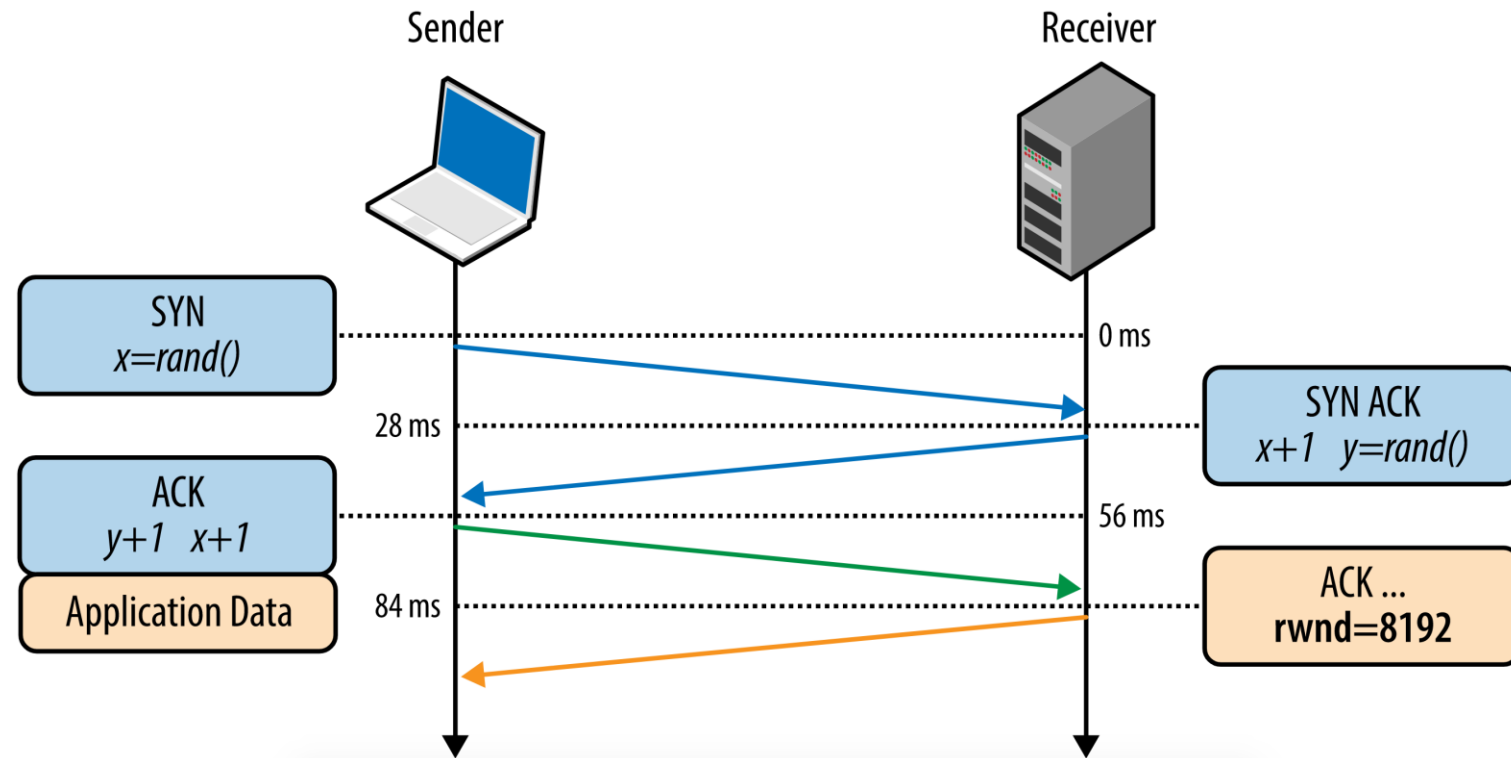
Congestion Avoidance and Control

- TCP has three mechanisms for preventing an application from overflowing the network.
 - Flow Control
 - Congestion Avoidance
 - Congestion Control
- Governs the rate at which data is placed onto the network.
- Which means, TCP governs the rate at which an application sends data.

Flow Control

- *Prevent the sender from overwhelming the receiver with data it cannot process.*
- Done through advertising a **window** of availability


- How many bytes are available to be recv'd
- Both Sender/Receiver will advertise these windows




(sort of) Slow Start

- Start slow, find out where the network starts to break down (quickly), hang out around that area window size
- Turns out, maybe not quick enough.
- HTTP tends to be “*bursty*”: many small-ish files needed to complete a web application.

- Every new connection needs to start this process anew.



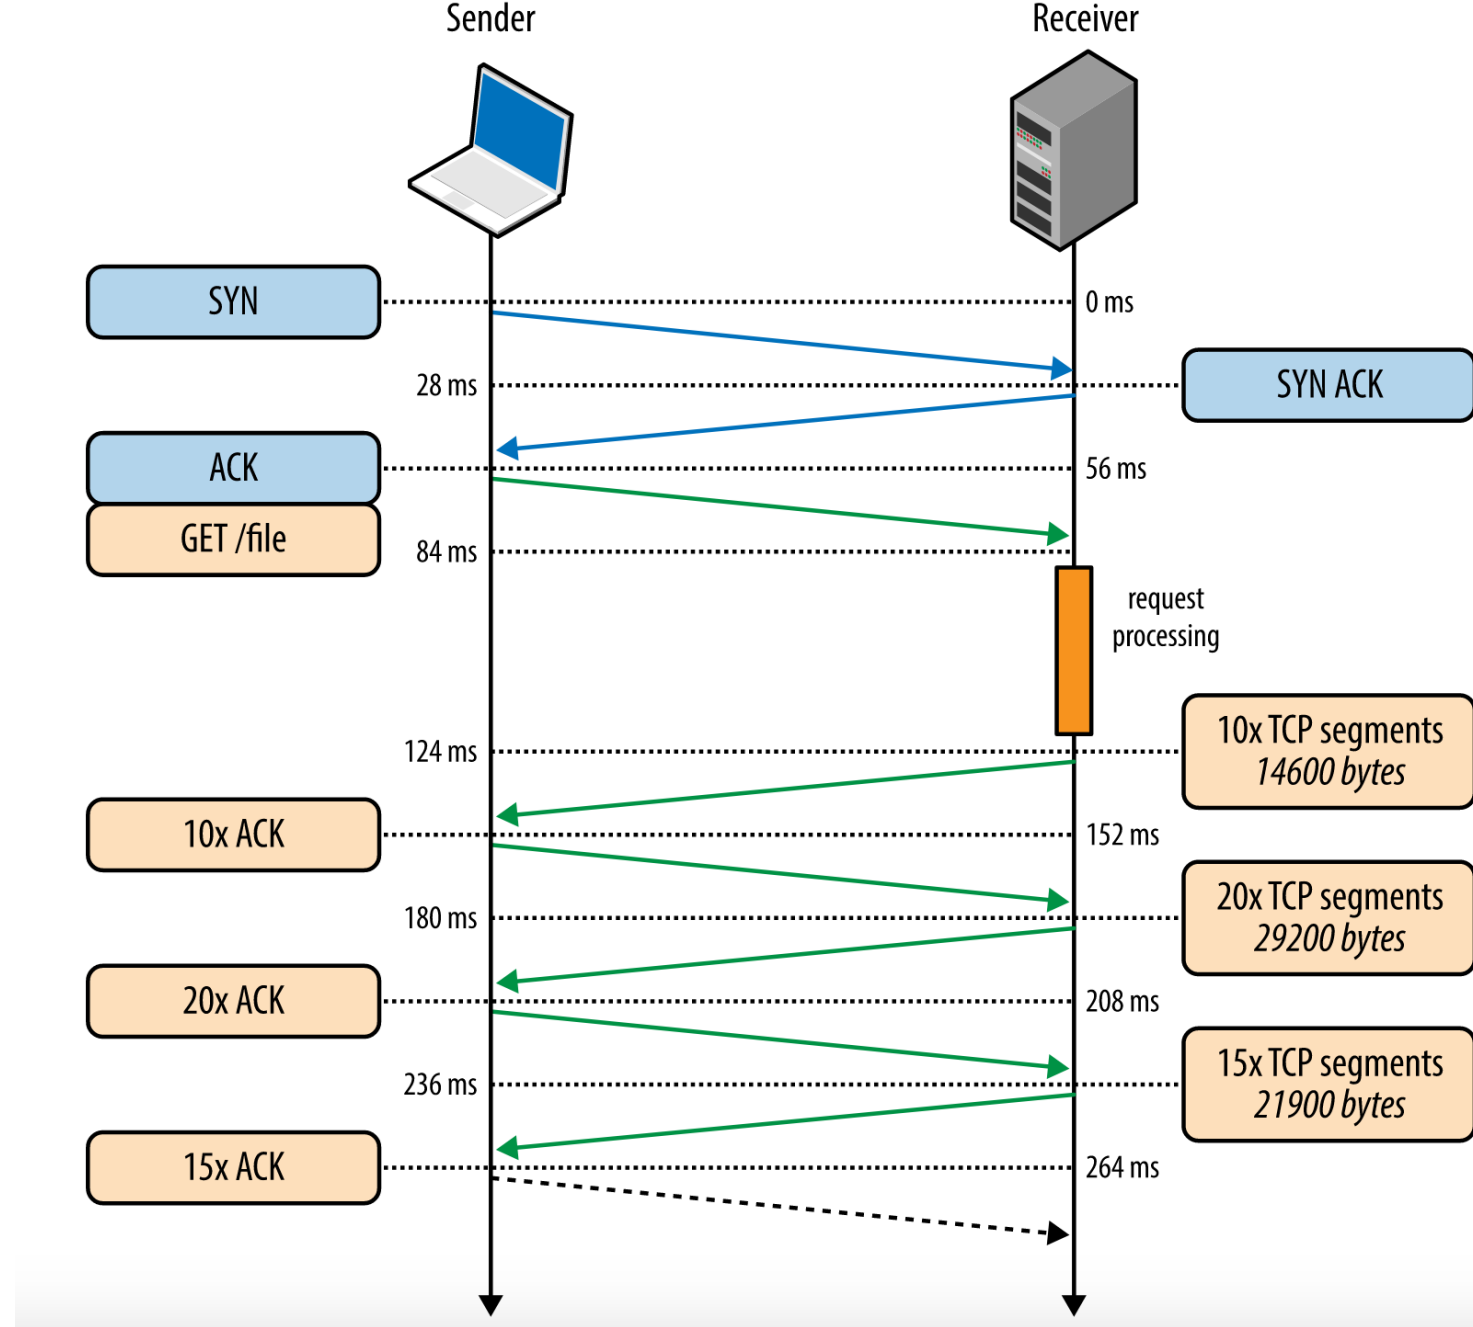
Congestion window size doubling every RTT



In a world where....


- RTT \approx 56 ms
- Client/Server Bandwidth: 5mbps
- Client/Server recv window: 64k bytes
- Initial Congestion Window: 10 segments (1460 bytes per segment)
 - 14k bytes per 10 segments
- Server processing time \approx 40ms
- Assume:
 - No packet loss
 - ACK per packet
 - GET request fits in single segment
- *Up next: Server wants to respond with 64k bytes of data. So what happens?*

- **0 ms** Client begins the TCP handshake with the SYN packet.
- **28 ms** Server replies with SYN-ACK and specifies its rwnd size.
- **56 ms** Client ACKs the SYN-ACK, specifies its rwnd size, and immediately sends the HTTP GET request.
- **84 ms** Server receives the HTTP request.
- **124 ms** Server completes generating the 64 KB response and sends 10 TCP segments before pausing for an ACK (initial cwnd size is 10).
- **152 ms** Client receives 10 TCP segments and ACKs each one.
- **180 ms** Server increments its cwnd for each ACK and sends 20 TCP segments.
- **208 ms** Client receives 20 TCP segments and ACKs each one.
- **236 ms** Server increments its cwnd for each ACK and sends remaining 15 TCP segments.
- **264 ms** Client receives 15 TCP segments, ACKs each one.





and if we reused that connection for another 64k file

- **0 ms** Client sends the HTTP request.
- **28 ms** Server receives the HTTP request.
- **68 ms** Server completes generating the 64 KB response, but the cwnd value is already greater than the 45 segments required to send the file; hence it dispatches all the segments in one burst.
- **96 ms** Client receives all 45 segments, ACKs each one.



Bandwidth vs. Latency



Lesson

- Bandwidth was irrelevant here (especially since we live in the future)
 - Latency is royalty (updated from Latency is king, then again...down with the monarch!)
 - Latency is the thing that matters the most!
 - The time it takes for a packet to get from one place to another
- Same request
 - 275% improvement in performance
- Lesson:
 - *Reusing connections has huge payoffs*
 - *Why use multiple connections for sending streams of data?*

Side note: linux kernel

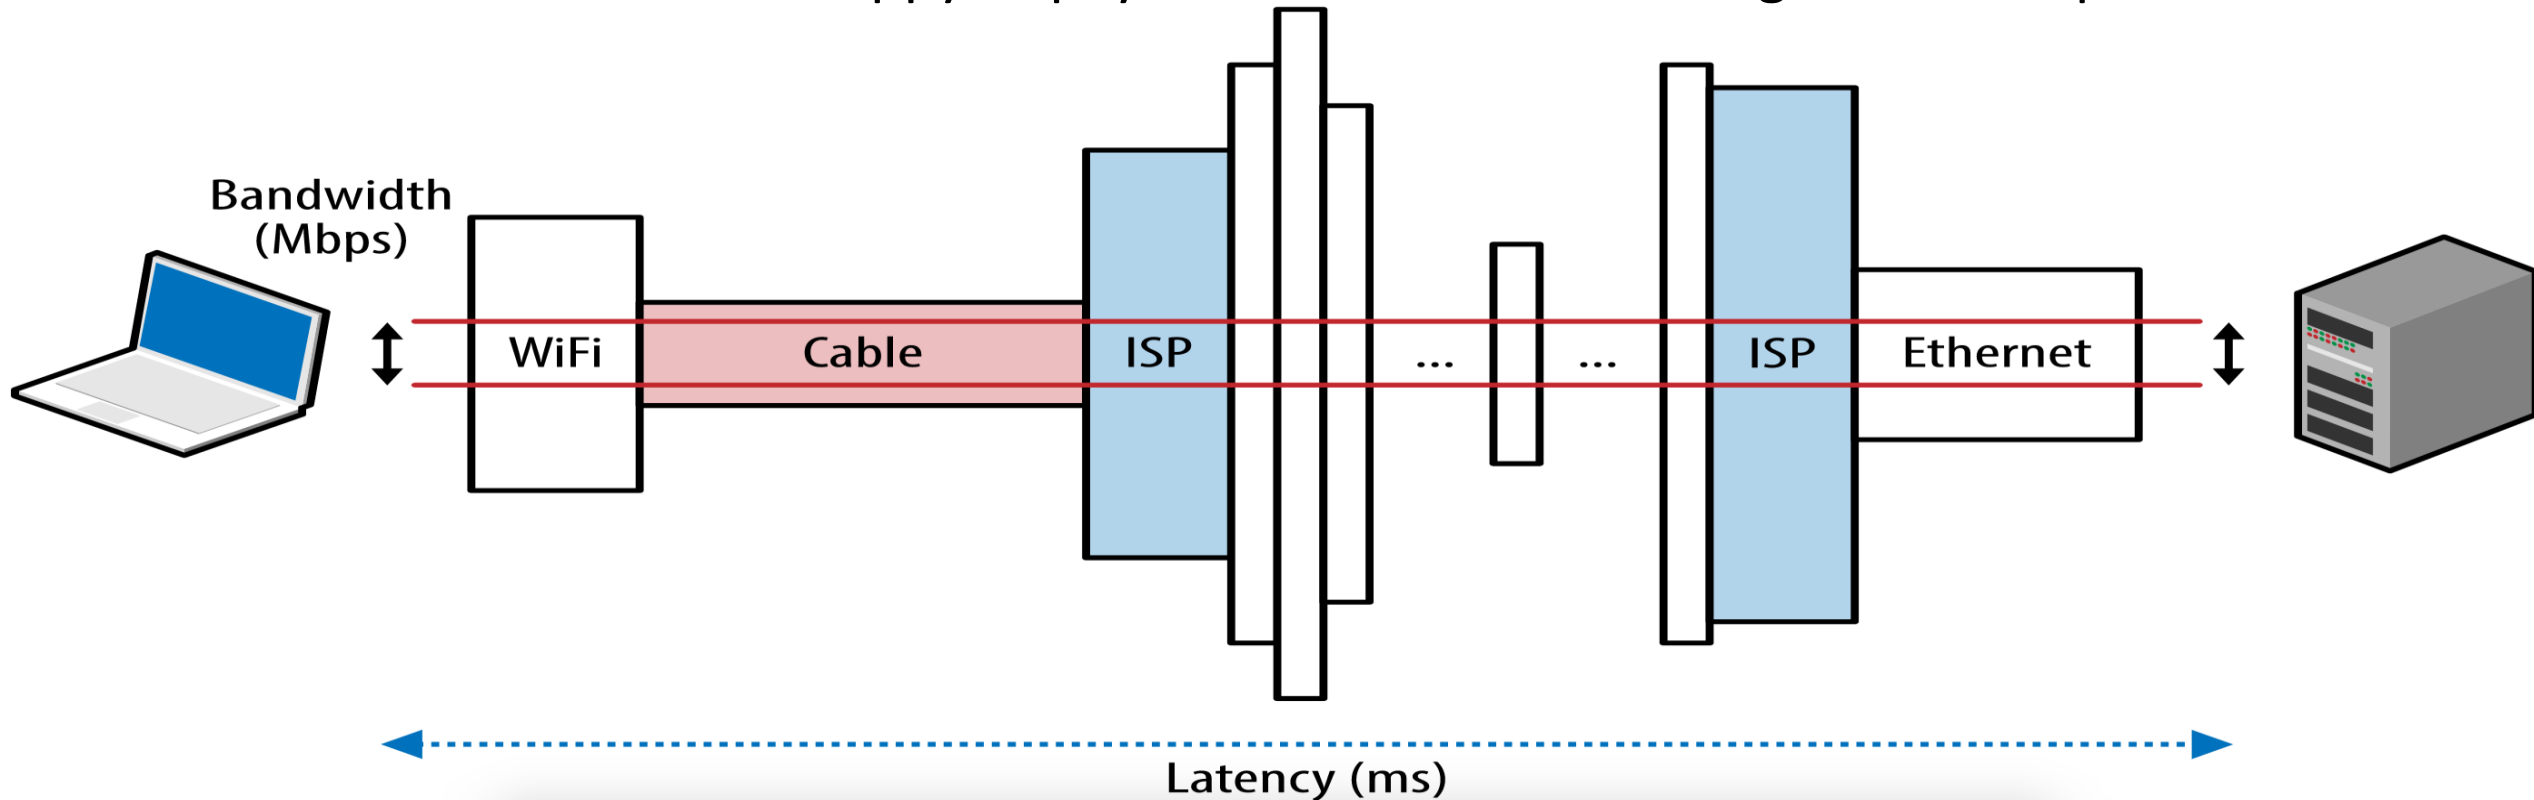
- You can view various TCP settings the kernel has with the command:
- `$ sysctl -a | egrep net.ipv4`
- `sysctl net.ipv4.tcp_slow_start_after_idle`
 - TCP will reset slow start if a connection has been idle
 - This setting can be disabled
- `sysctl net.ipv4.tcp_window_scaling`
 - Check if window scaling is enabled.
- Do not blindly modify kernel settings... ask an expert
- Keep your kernels updated.

Congestion Avoidance

- *Packet loss **will** happen*
 - Somewhere along the way we encounter a router, or some other node that is overwhelmed.
 - The connection must the adjust the window to avoid further packet loss
- Congestion Avoidances kicks in after slow start to grow the window but avoid further loss.
 - After exponential, small increases to the window size
- Still an active area of research
 - Kernel changes
 - Tweaks available
 - Money maker

onto bandwidth and latency

- Speed costs money
 - 2015, Hibernia spent 300+ million dollars to create a new route between New York at London
 - ~5 ms faster than another other transatlantic link
 - **\$60,000,000 per millisecond saved (\$60,000,000 for 1ms smaller ping)**
 - Financial markets happy to pay this to have an advantage over competitors



Why we care about delay

General rule of thumb is a web page should be displayed **within 250ms** otherwise the user may feel something is off.

Delay	User perception
0–100 ms	Instant
100–300 ms	Small perceptible delay
300–1000 ms	Machine is working
1,000+ ms	Likely mental context switch
10,000+ ms	Task is abandoned

HTTP/2

- **Goals** (incomplete list)
 - **Reduce latency** with full multiplexing
 - **Minimize overhead** with efficient compression of HTTP headers
 - Add support for prioritization and server push
- **HTTP semantics are not modified**
 - Same HTTP methods, status codes, URIs, headers, etc.
 - Smarter Data Formatting instead
 - Result: any application can be upgraded from 1.1 -> 2 without modification
- **Performance in mind**
 - We can still improve performance

Timeline

- March 2012: Call for proposals for HTTP/2
- November 2012: First draft of HTTP/2 (based on [SPDY](#))
- August 2014: HTTP/2 draft-17 and HPACK draft-12 are published
- August 2014: Working Group last call for HTTP/2
- February 2015: IESG approved HTTP/2 and HPACK drafts
- May 2015: RFC 7540 (HTTP/2) and RFC 7541 (HPACK) are published

SPDY developed at Google

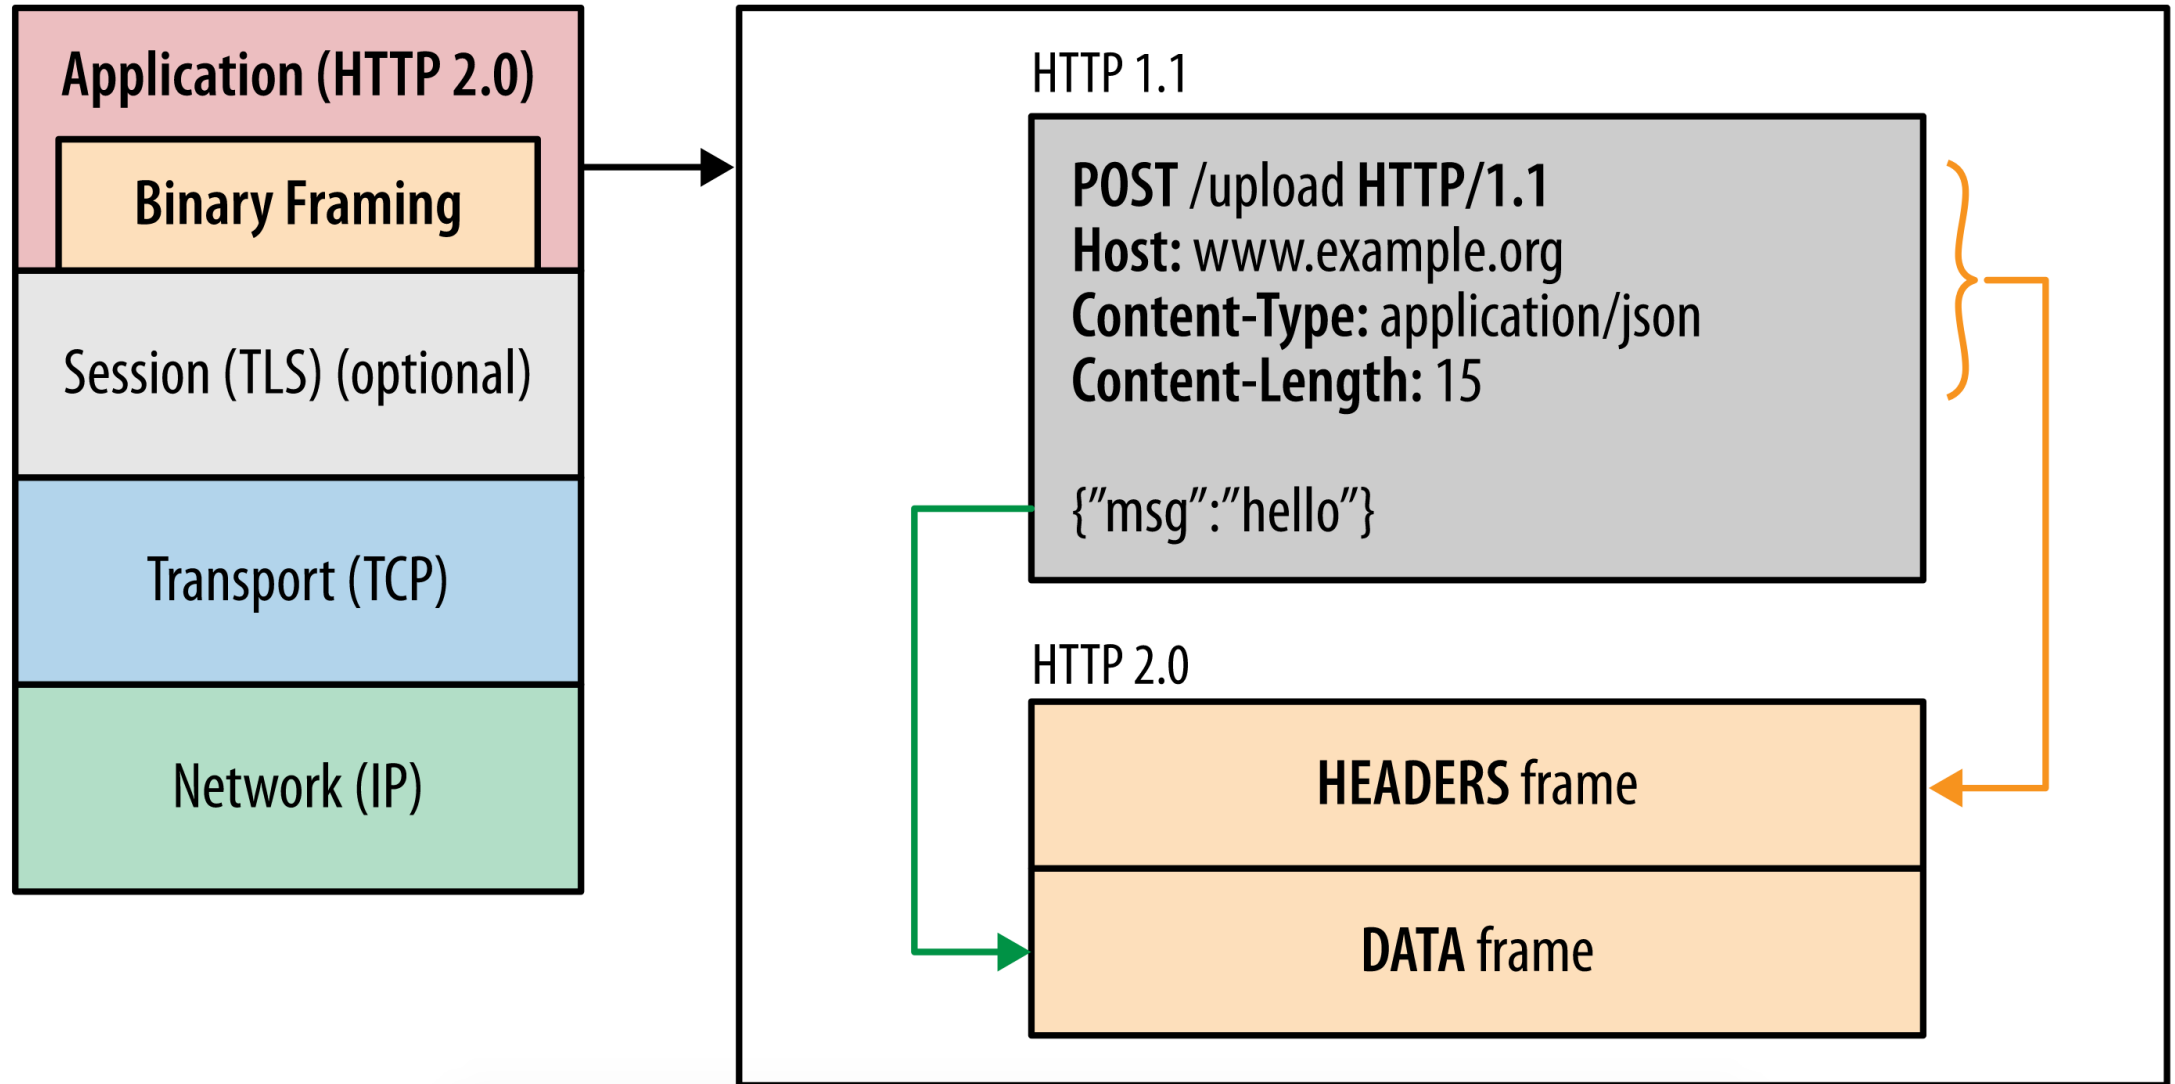
Drama

- Relatively short timeline, did others have a chance to implement an alternative?
 - Microsoft alternative:
 - https://en.wikipedia.org/wiki/HTTP_Speed%2BMobility
 - Facebook's response: <https://lists.w3.org/Archives/Public/ietf-http-wg/2012JulSep/0251.html>
 - Network-Friendly HTTP Upgrade (alternative): <https://tools.ietf.org/html/draft-tarreau-httpbis-network-friendly-00>
 - FB's response contains logical and sound reasoning for not considering this alternative
- Encryption as a Requirement?
 - CPU requirements
 - Which encryption protocol?
 - Working group decided against it, though all clients who implement HTTP/2 have made it mandatory in the end.
 - Makes our HTTP/2 captures a little more painful since we have to break encryption

Drama

- BINARY PROTOCOL?!?!?
 - HTTP/2 is no longer an ASCII Protocol
 - Meaning, you need a tool to inspect HTTP/2 packets (like Wireshark)
 - This bothers some people I guess
- HTTP/3 is coming
 - HTTP/2 + TLS + QUIC
 - Switching from TCP -> UDP (sort of)
 - Discussed later in the course

Biggest Change: Binary Framing Layer



ASCII Protocol (HTTP/1.1)

▼ Hypertext Transfer Protocol

► GET / HTTP/1.1\r\n

Host: scrivnor.cikeys.com\r\n

User-Agent: curl/7.64.1\r\n

Accept: */*\r\n

Connection: Upgrade, HTTP2-Settings\r\n

Upgrade: h2c\r\n

HTTP2-Settings: AAMAAABkAARAAAAAAIAAAAA\r\n

\r\n

[\[Full request URI: http://scrivnor.cikeys.com/\]](http://scrivnor.cikeys.com/)

[HTTP request 1/1]

0030	08 0a fc ed 00 00 01 01	08 0a 51 c1 6c 6c 02 56Q.ll.V
0040	82 67 47 45 54 20 2f 20	48 54 54 50 2f 31 2e 31	·gGET / HTTP/1.1
0050	0d 0a 48 6f 73 74 3a 20	73 63 72 69 76 6e 6f 72	··Host: scrivnor
0060	2e 63 69 6b 65 79 73 2e	63 6f 6d 0d 0a 55 73 65	.cikeys. com··Use
0070	72 2d 41 67 65 6e 74 3a	20 63 75 72 6c 2f 37 2e	r-Agent: curl/7.
0080	36 34 2e 31 0d 0a 41 63	63 65 70 74 3a 20 2a 2f	64.1··Ac cept: */
0090	2a 0d 0a 43 6f 6e 6e 65	63 74 69 6f 6e 3a 20 55	*··Conne ction: U
00a0	70 67 72 61 64 65 2c 20	48 54 54 50 32 2d 53 65	pgrade, HTTP2-Se
00b0	74 74 69 6e 67 73 0d 0a	55 70 67 72 61 64 65 3a	ttings·· Upgrade:
00c0	20 68 32 63 0d 0a 48 54	54 50 32 2d 53 65 74 74	h2c··HT TP2-Sett
00d0	69 6e 67 73 3a 20 41 41	4d 41 41 41 42 6b 41 41	ings: AA MAAABkAA
00e0	52 41 41 41 41 41 41 41	49 41 41 41 41 41 0d 0a	RAAAAAAA IAAAAA··

Binary Protocol (HTTP/2)

- HTTP Headers and Data are encoded in a binary format (not ASCII)

▼ Stream: HEADERS, Stream ID: 1, Length 77, 200 OK

Length: 77

Type: HEADERS (1)

► Flags: 0x04

0... .. = Reserved: 0x0

.000 0000 0000 0000 0000 0000 0000 0001 = Stream Identifier: 1

[Pad Length: 0]

Header Block Fragment: 886196df3dbf4a002a651d4a05f740a0017000b800298b46...

[Header Length: 210]

[Header Count: 7]

► Header: :status: 200 OK

► Header: date: Thu, 01 Jan 1970 00:00:00 GMT

0040	6c 6c 00 00 4d 01 04 00	00 00 01 88 61 96 df 3d	ll..M... ..a..=
0050	bf 4a 00 2a 65 1d 4a 05	f7 40 a0 01 70 00 b8 00	.J.*e.J. .@..p...
0060	29 8b 46 ff 76 85 86 b1	92 72 ff 6c 96 d0 7a be)·F·v... ·r·l...z·
0070	94 13 aa 65 1d 4a 08 02	02 80 6e e3 42 b8 27 54	...e·J.. ··n·B·'T
0080	c5 a3 7f 52 84 8f d2 4a	8f 0f 0d 83 6d c7 df 5f	...R...J ...m..._
0090	87 49 7c a5 89 d3 4d 1f	00 16 43 00 01 00 00 00	·I ...M· ..C.....

Key Terms

- **Stream**

- Bi-directional flow of data, may carry one or more **messages**

- **Message**

- A sequence of **frames** that map to a logical request/response.

- **Frame**

- Smallest unit of communication.
- Carries a specific type of data, headers, message payload, etc.

- A **stream** is a bi-directional flow of **frames**. Those **frames** are received and put back together to form the **message**.

Stream 1

Request message

HEADERS frame (stream 1)

:method: GET
:path: /index.html
:version: HTTP/2.0
:scheme: https
user-agent: Chrome/26.0.1410.65

Response message

HEADERS frame (stream 1)

:status: 200
:version: HTTP/2.0
server: nginx/1.0.11
vary: Accept-Encoding
...

DATA frame (stream 1)

... response payload...

Stream N

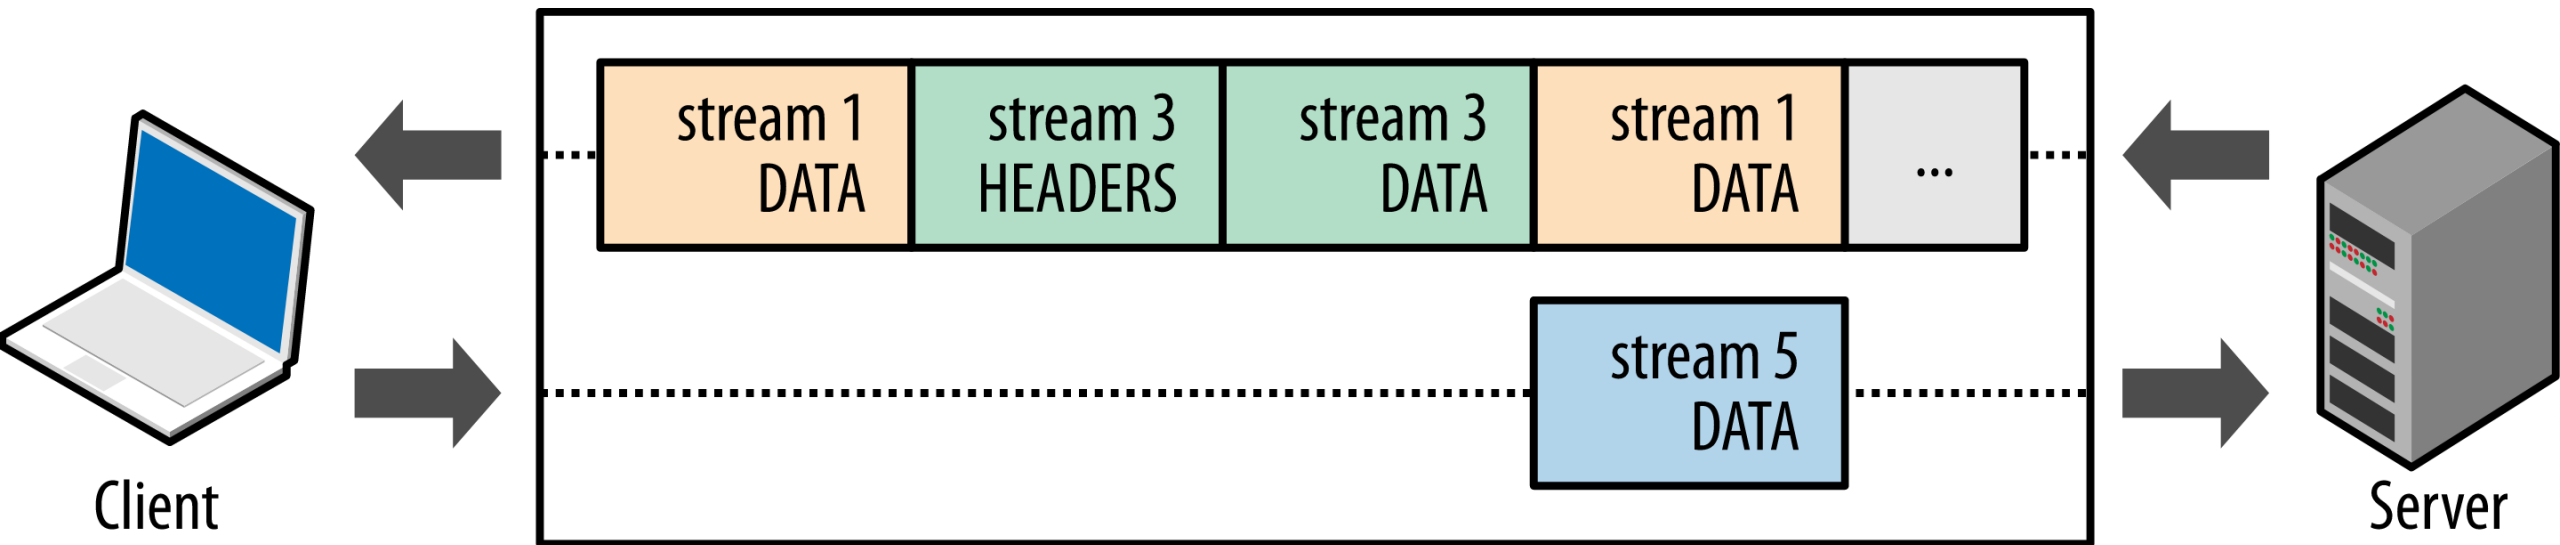
⋮



Simple change, many benefits


- By adding the ability to split messages up into frames that can be interleaved over a stream, we get
 - **Interleave** multiple requests in parallel *without blocking* on any one
 - **Interleave** multiple responses in parallel *without blocking* on any one
 - **Use a single connection to deliver multiple requests and responses in parallel**
 - Remove unnecessary HTTP/1.x workarounds
 - Deliver lower page load times by eliminating unnecessary latency and improving utilization

HTTP 2.0 connection




Stream Prioritization

- Recall, web applications have dependencies and order matters
- Stream **dependencies** (one stream depends on another)
- Stream **priorities** (one message has a higher priority over another)

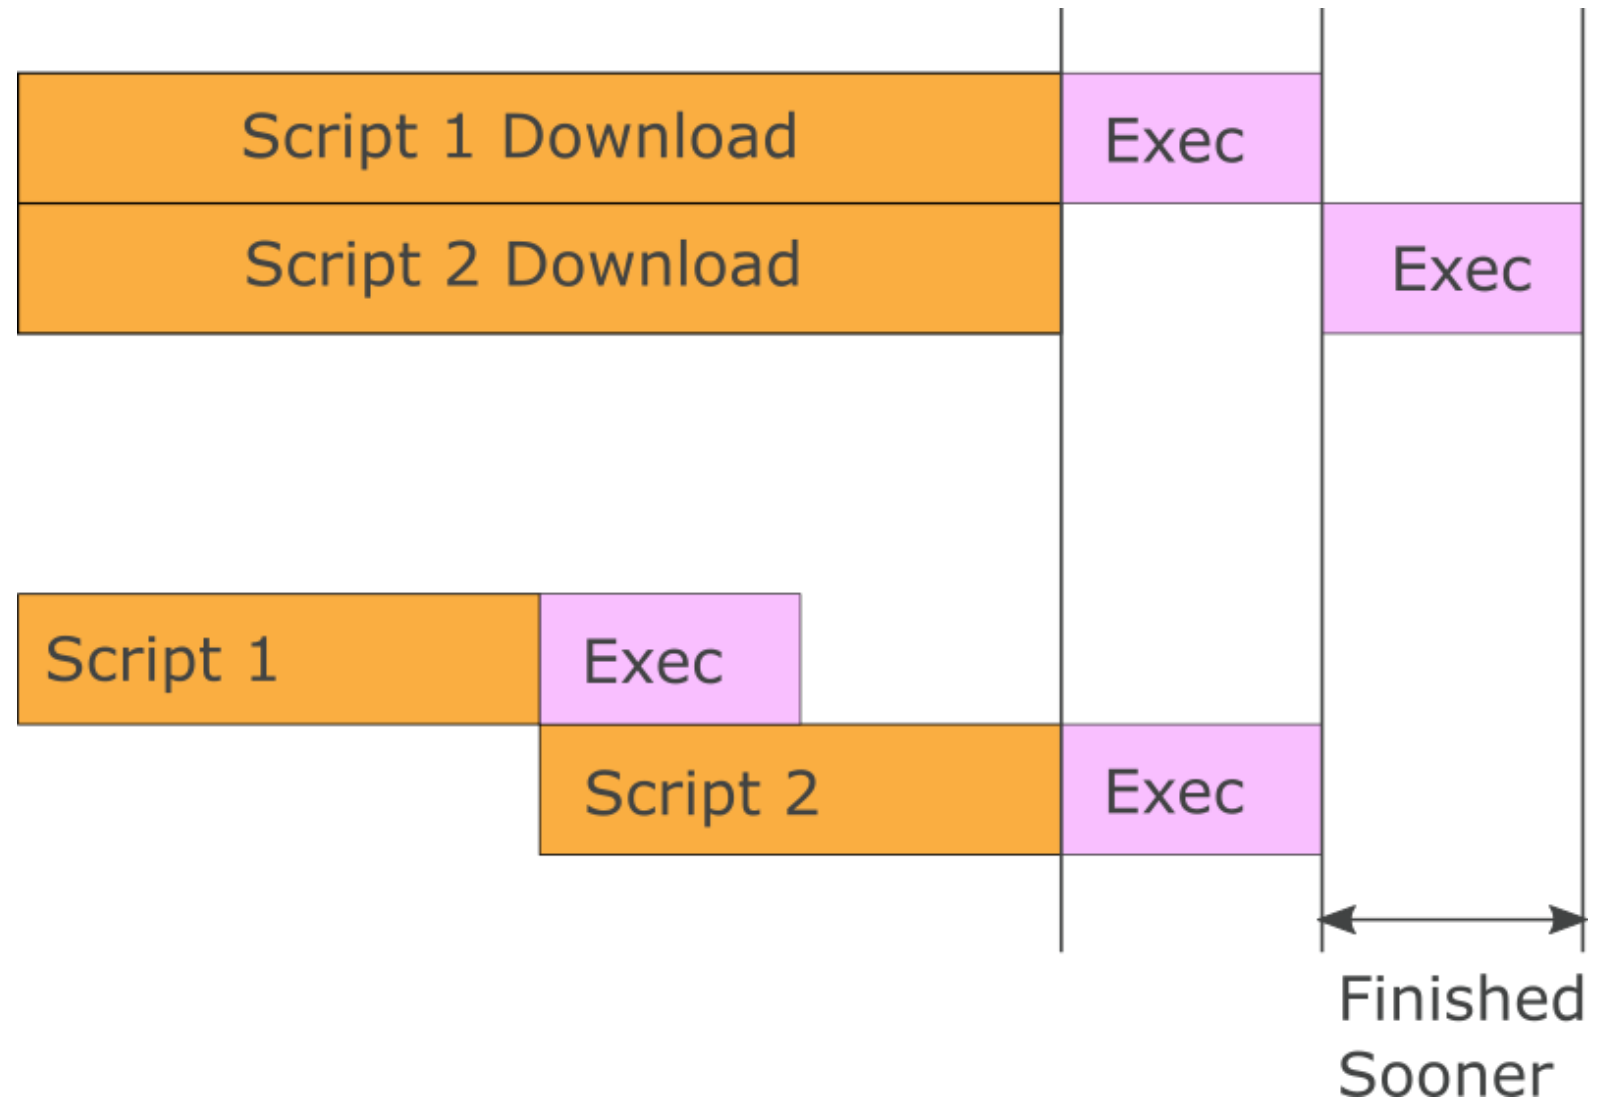


Example: Prioritization Tree

- A+B
 - Same dependency level in the tree
 - Total weight: $W(A) + W(B) = 12 + 4 = 16$
 - A then gets 12/16 of the bandwidth
 - B then gets 4/16 of the bandwidth
- D+C
 - C depends on D, therefore D gets full bandwidth until finished.
- Weights are simply integers 1-255
- Note: a client may update its priority *preference*, but it can't tell the server what to do.
- <https://blog.cloudflare.com/better-http-2-prioritization-for-a-faster-web/>



How does it help?



Just One TCP Connection (per origin)

- HTTP/2 only requires a **single TCP connection**
- **Benefits**
 - TCP designed for long-lived connections, HTTP is short/bursty
 - Fewer TLS handshakes
 - Overall reduction of resource demand on servers/clients
- TCP's head of line blocking **not solved**
 - During packet loss, TCP will hold onto data while trying to recover the missing
 - **HTTP head of line blocking is solved by HTTP/2**
- Packet loss reduces maximum throughput for **entire connection** now, not just one of twelve, say
- A well configured TCP on a server can lead to huge gainz
- TCP not the only choice, the **future may be QUIC**

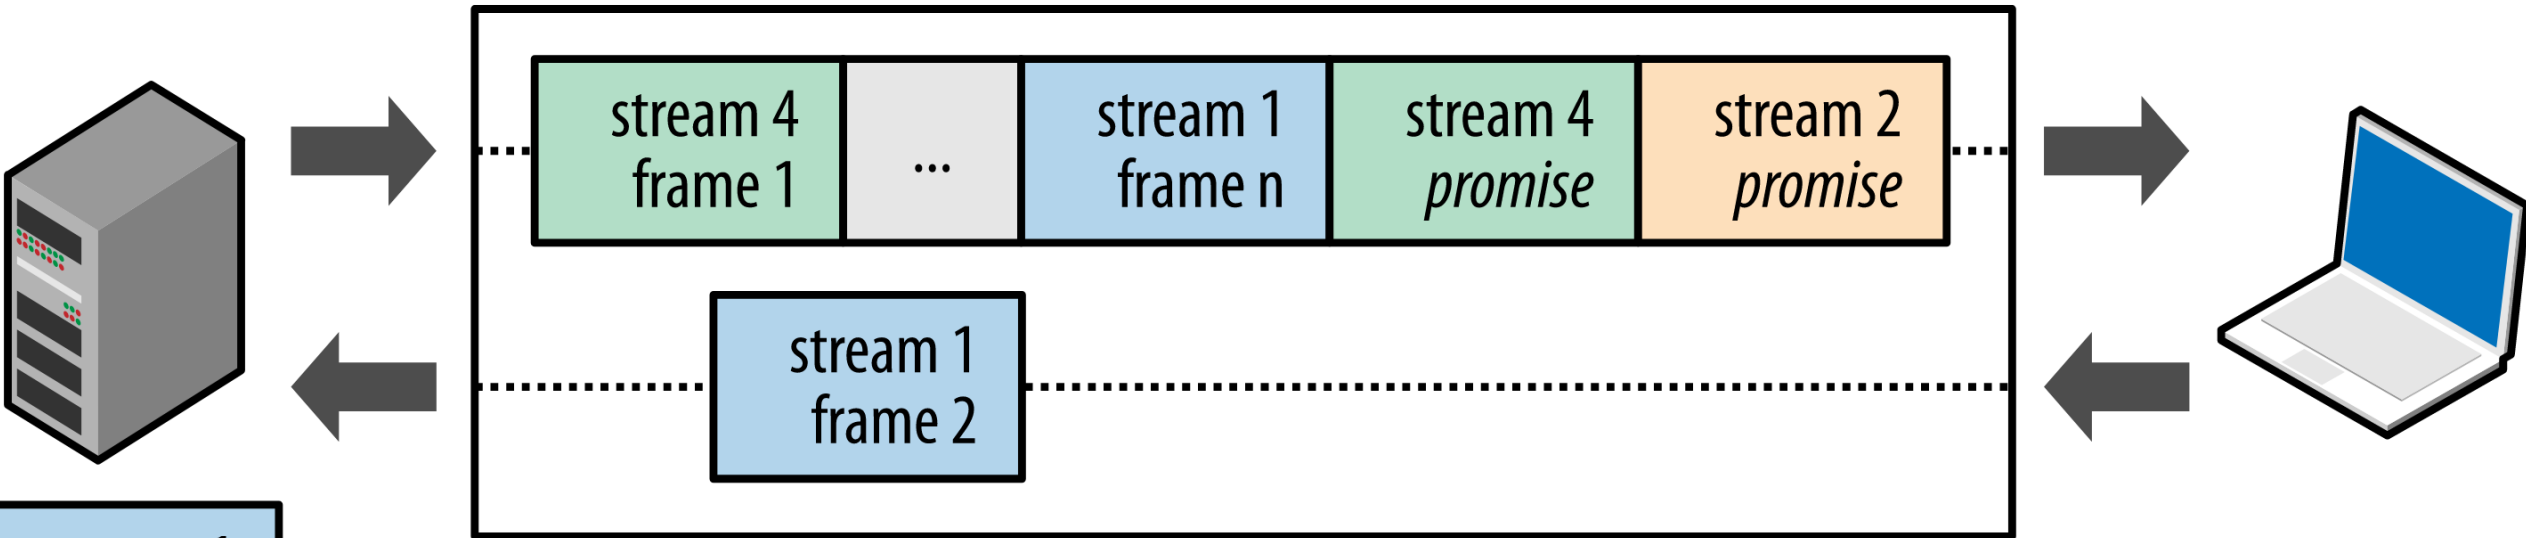
HTTP/2 Flow Control

- Slightly controversial because the transport layer has flow control already.
 - Reimplementing features at a higher layer?
- The concept is roughly the same as TCP, so we will study this concept later in the course.
- For now, just know:
 - Directional
 - Credit based
 - Can be updated

Server Push

- Client receives a promise, can accept or deny it

HTTP 2.0 connection



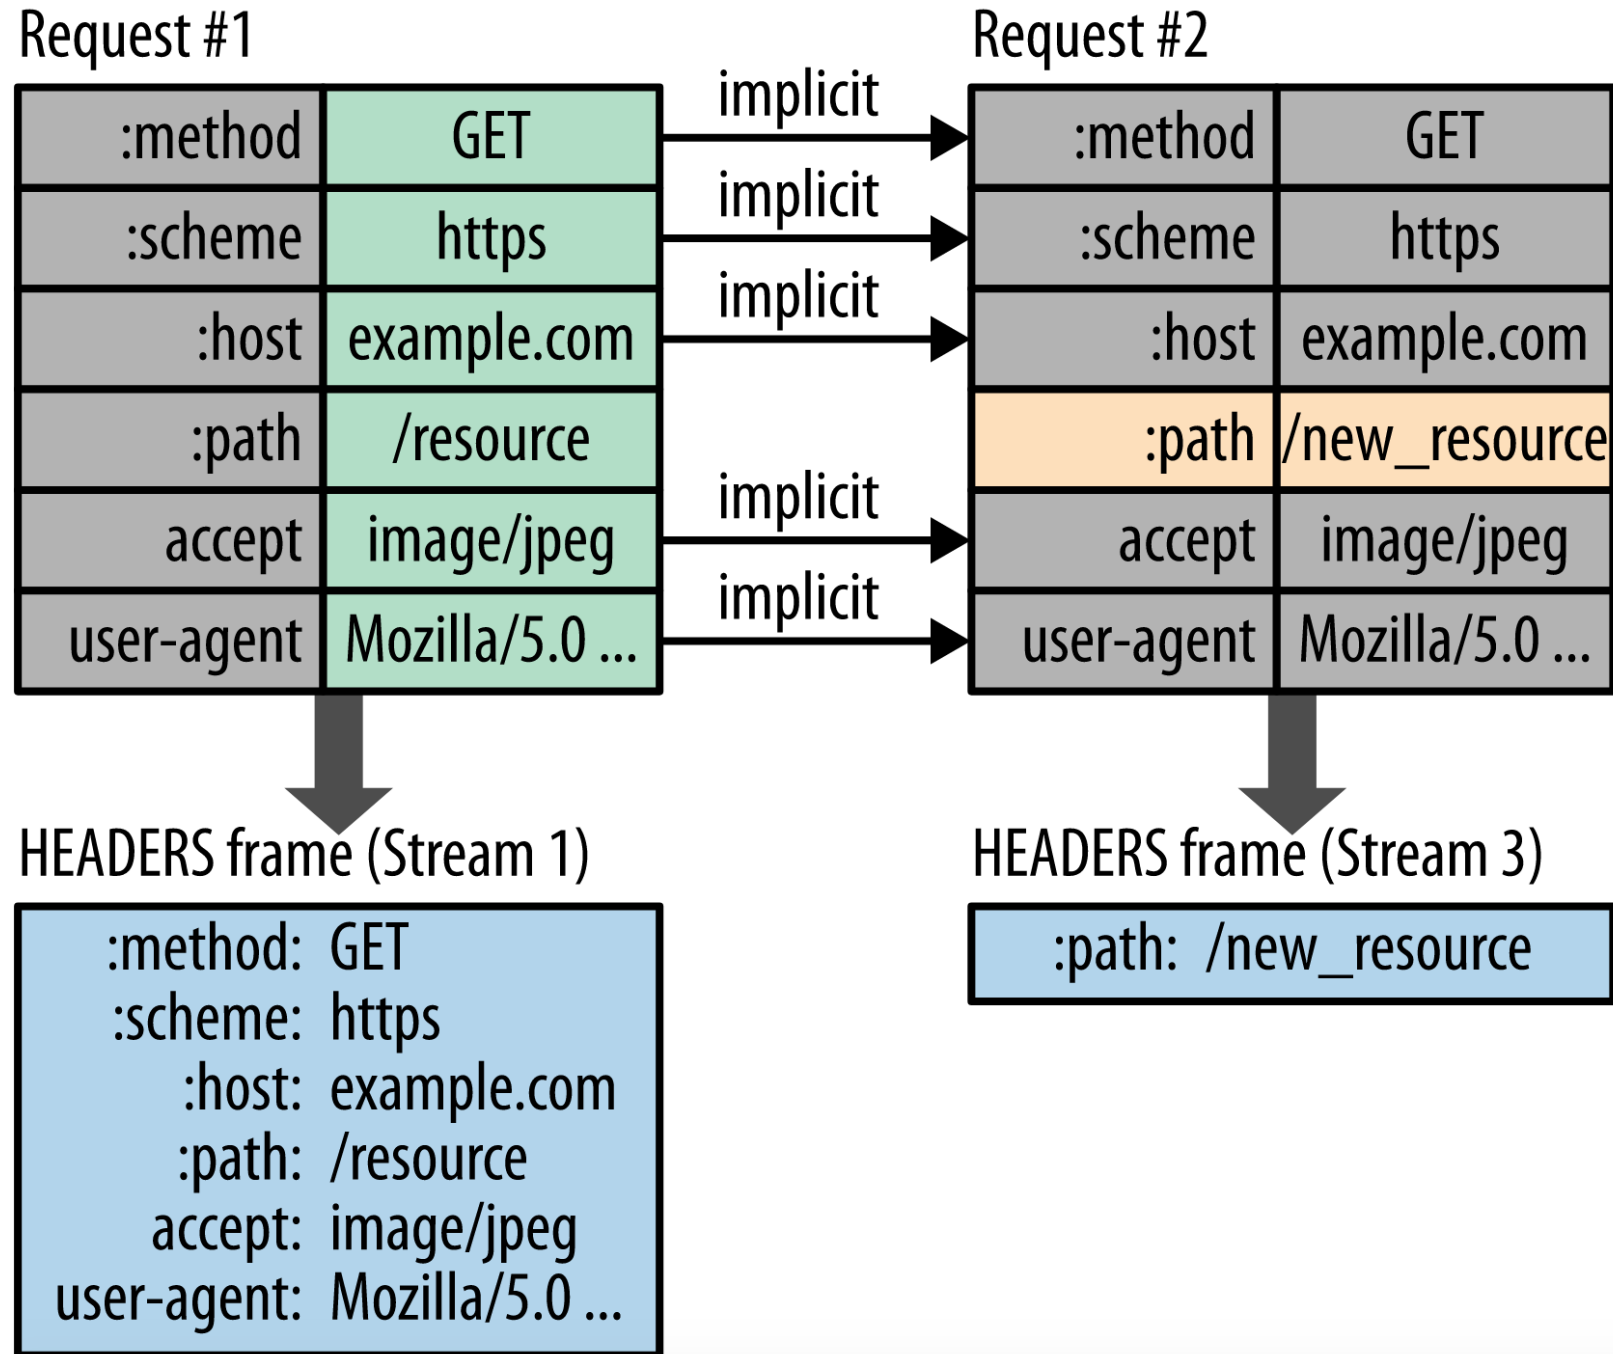
stream 1: /page.html (client request)

stream 2: /script.js (push promise)

stream 4: /style.css (push promise)

Optimizing headers

- Huffman encoding
- Implicit headers have no need to be sent, redundant
- Headers can add 500-800 bytes of data




► HyperText Transfer Protocol 2


0000	8c	85	90	31	3c	75	fa	5b	3b	20	40	6d	08	00	45	00	...	1<u	[;	@m	...	E	.				
0010	00	45	6d	55	40	00	33	06	cd	52	8a	c5	c0	4e	c0	a8	...	EmU@	.	3	.	R	...	N	...			
0020	01	4f	00	50	ca	2d	39	9b	ff	85	83	37	af	2a	80	18	...	0	.	P	.	-9	7	.	*	...
0030	00	eb	ad	fd	00	00	01	01	08	0a	02	be	85	78	52	28	xR	(
0040	a1	49	00	00	08	07	00	00	00	00	00	00	00	00	01	00	...	I		
0050	00	00	00																							

Bit	+0..7	+8..15	+16..23	+24..31
0	Length			Type
32	Flags			
40	R	Stream Identifier		
...	Frame Payload			


Looking at HTTP/1.1 Traffic



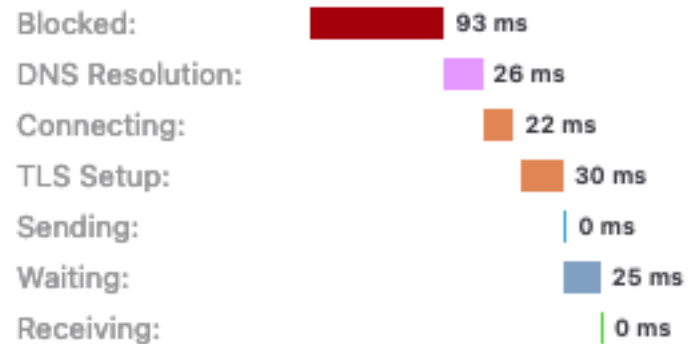
Request Timing






Compared to HTTP/2.0




Request Timing




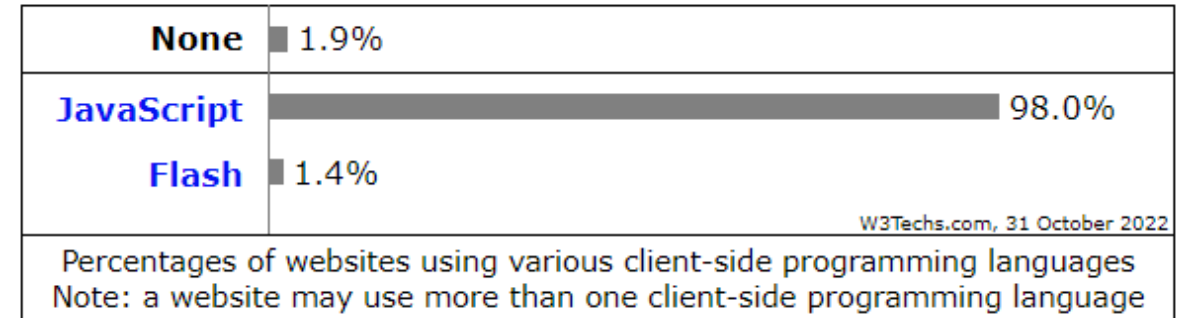

Web Statistics: web servers




Web Statistics: elements



Web Statistics: powered by...



Web Statistics: advertising



HTTP/3 is ~~coming soon~~ here!

- <https://datatracker.ietf.org/doc/draft-ietf-quic-http/>
- Version 30 (Fall 2020)
- Version 34 (Spring 2021)
- RFC 9114 (Summer 22)
- HTTP/2 + QUIC + TLS
- Implemented in most major browsers at this point.
- Major changes:
 - Use UDP instead of TCP
 - Encrypted traffic required




JSON Review

- JavaScript Object Notation
- Extension: `.json`
- Language independent format derived from JavaScript.
- Pronounced: jay-sawn
 - Creator says it is pronounced “Jason” but also, “doesn’t care”
- Early 2000’s – present

JSON Example: course schedule

```
{  
  "course": "COMP 221",  
  "semester": "Fall 2021",  
  "title": "UNIX Systems Programming I",  
  "schedule": {  
    "headers": [  
      "Week", "Topics", "Reading", "Assignments"  
    ],  
    "weeks": [  
      { "type": "week", "data": {  
        "topic": "Getting Started",  
        "reading": "Ch 2, Ch 3",  
        "assignments": "L01"  
      }  
    }  
  ]  
}
```

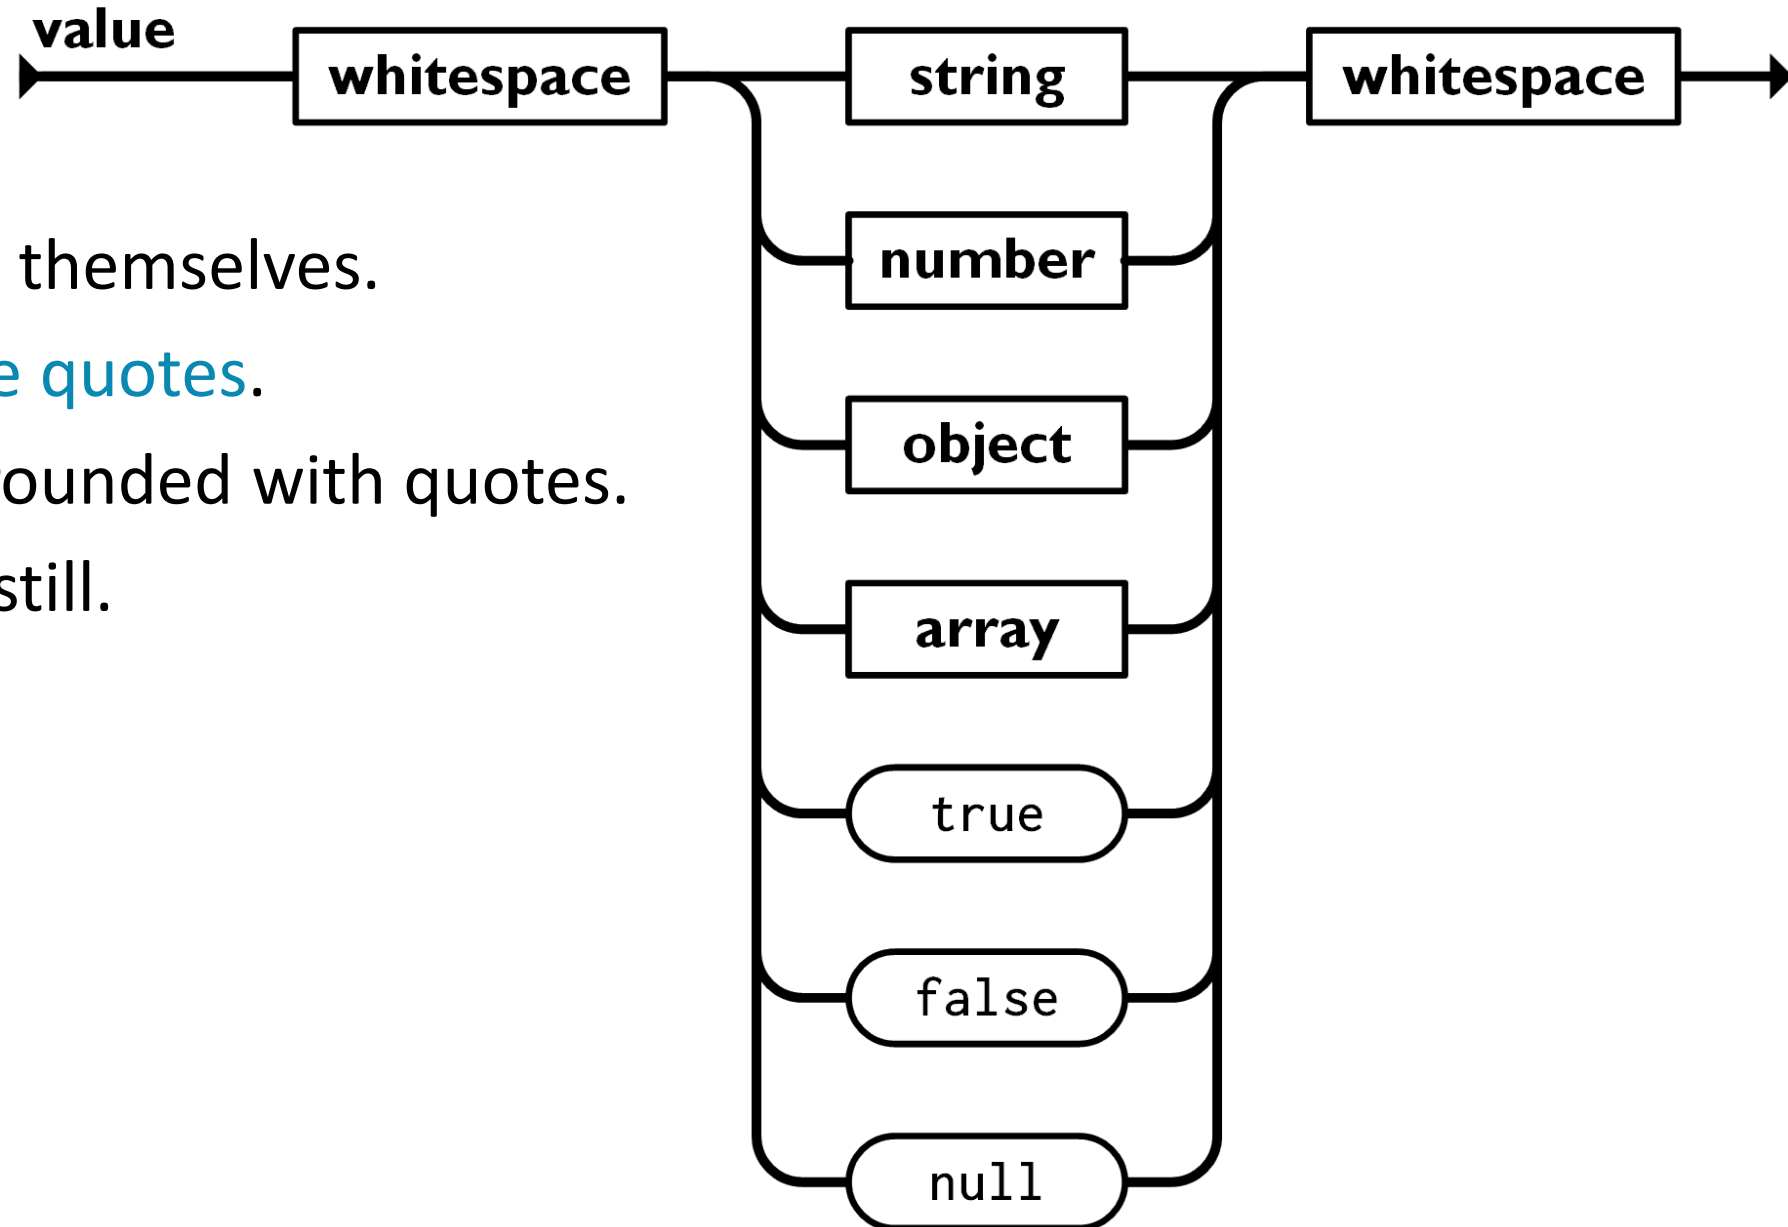
JSON: Object specification



Need to define
string and value still

```
{ "course": "COMP 221", "semester": "Fall 2021" }
{ }
{ "course": "COMP 221" }
```

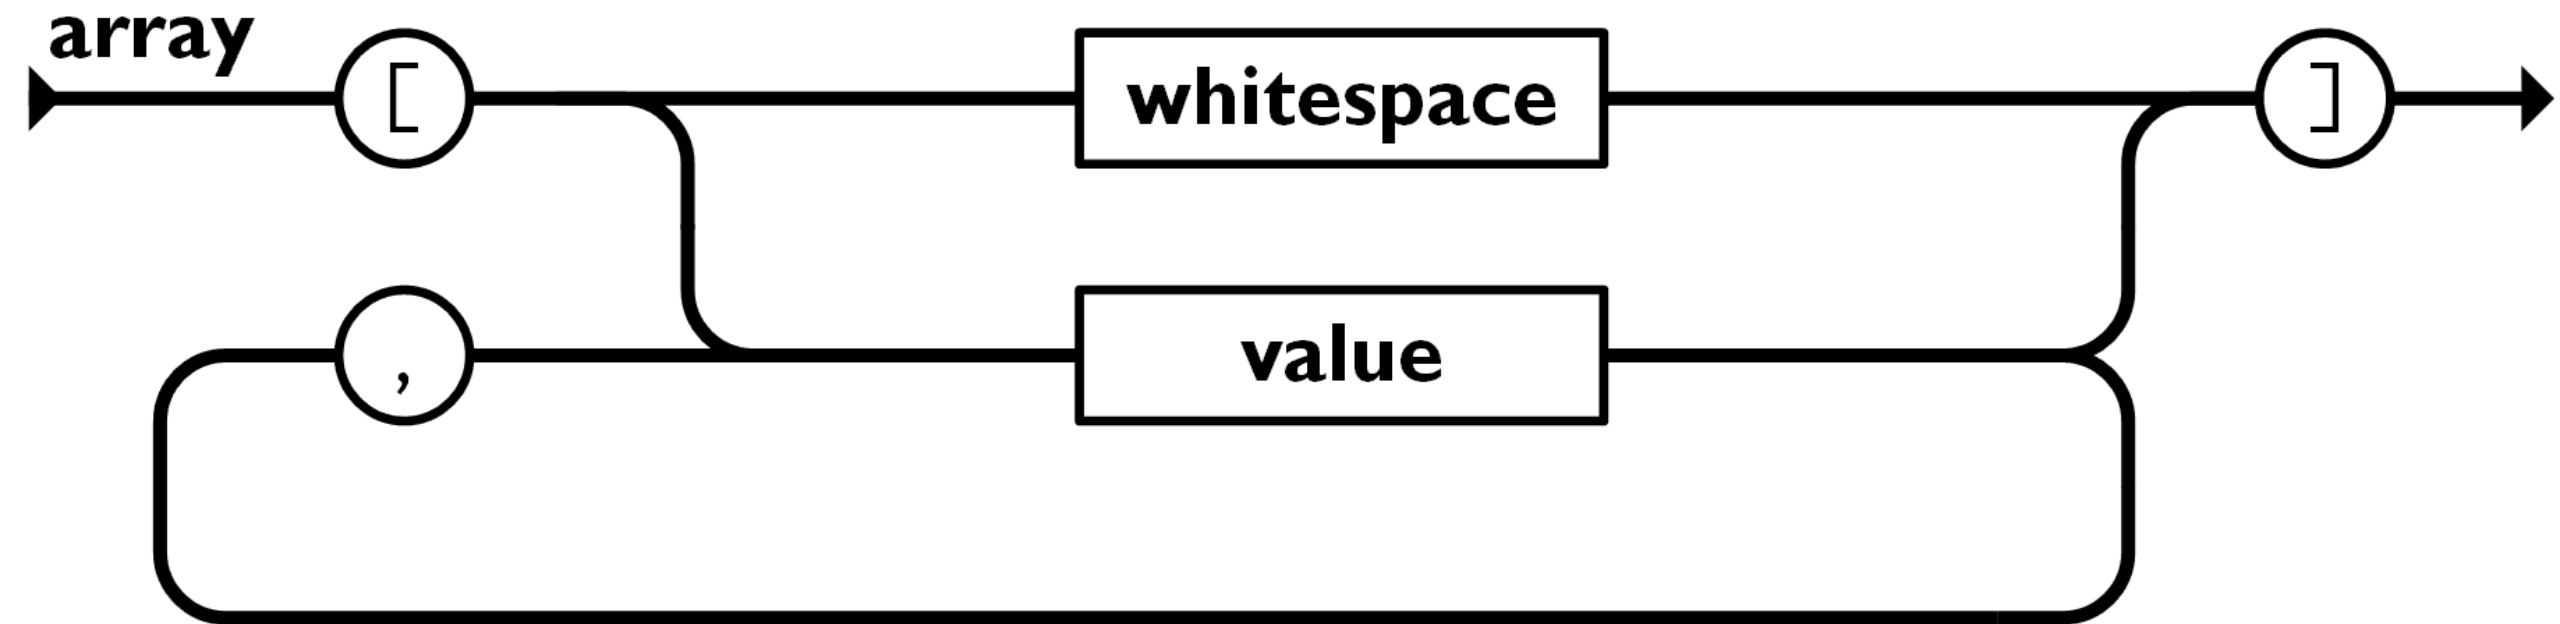
JSON: Value Specification



- Values can be **objects** themselves.
- Strings require **double quotes**.
- **Numbers** are not surrounded with quotes.
- Need to define array still.

JSON: Array specification

- Comma separated values surrounded with [].



HAR (Http ARchive) files are JSON

- Example from documentation
- <https://w3c.github.io/web-performance/specs/HAR/Overview.html>

request

This object contains detailed info about performed request.

```
"request": {  
  "method": "GET",  
  "url": "http://www.example.com/path/?param=value",  
  "httpVersion": "HTTP/1.1",  
  "cookies": [],  
  "headers": [],  
  "queryString" : [],  
  "postData" : {},  
  "headersSize" : 150,  
  "bodySize" : 0,  
  "comment" : ""  
}
```